

Copyright

by

Karan Prakash Hiranandani

2023

hIPPYfire:(yet-to-be-finalized)

APPROVED BY

SUPERVISING COMMITTEE:

Dr. Omar Ghattas, Supervisor

Dr. Umberto Villa, Co-Supervisor

hIPPYfire:(yet-to-be-finalized)

by

Karan Prakash Hiranandani

REPORT

Presented to the Faculty of the Graduate School
of the University of Texas at Austin
in Partial Fulfillment
of the Requirements
for the Degree of

MASTER OF SCIENCE

The University of Texas at Austin
May 2023

Acknowledgments

I would like to express my sincere gratitude to **Dr. Umberto Villa** for his supervision and guidance throughout this project. Not only were his insights and advice crucial in the conception of this report, but his teachings on this topic have been invaluable to me.

I am extremely grateful to **Dr. Omar Ghattas** for agreeing to supervise this project and assisting with its completion. I hope that the work done in this project proves to be useful to him in the research him and his group are conducting.

I would also like to thank **Stephanie Rodriguez**, who has ensured that everything leading to this point has been smooth sailing. And finally, I am indebted to all the faculty members, staff, colleagues, and peers at the Oden Institute—this journey would not have been possible without your guidance.

hIPPYfire:(yet-to-be-finalized)

by

Karan Prakash Hiranandani, M.S.

The University of Texas at Austin, 2023

Supervisor: Dr. Omar Ghattas

This study presents the implementation of **hIPPYfire**, a solver for large-scale Bayesian and deterministic inverse problems governed by partial differential equations (PDEs) with infinite-dimensional parameter fields that become high-dimensional after discretization. It utilizes the same scalable algorithms introduced by its predecessor, **hIPPYlib**, such as the inexact Newton Conjugate Gradient (Newton-CG) method for the computation of the maximum *a posteriori* distribution (MAP point) and the low rank-approximation of the Hessian. These algorithms exploit the fact that several PDE models of physical systems have a low-dimensional solution manifold. **hIPPYfire** computes the solution of the inverse problem at a cost independent of the parameter dimension, which is measured in terms of the number of linear forward PDE solves. However, unlike **hIPPYlib** (which is built on FEniCS), **hIPPYfire** uses Firedrake to solve the PDE governing the forward problem. Firedrake presents a unique modular structure that clearly distinguishes between the programming and mathematical aspects of the library—thereby enabling contributions from programmers and mathematicians alike and ensuring its consistent development. The functionality of the solver is validated by running it on an inverse problem that is governed by an elliptic PDE according to the Bayesian framework. The major components of the inverse problem, namely the forward problem, misfit, and prior functionals, are clearly defined and used to compute the solution of the forward problem and calculate the MAP point using the inexact Newton-CG method. The successful computation of the forward problem solution and MAP point, coupled with the efficient abstraction in Firedrake, provide motivation to incorporate functionality into **hIPPYfire**, such as the low-rank approximations of the posterior covariance and the Hessian of the data misfit.

Table of Contents

Acknowledgments	4
Abstract	5
List of Figures	7
Chapter 1. Intrduction	8
Chapter 2. Theory Behind Inverse Problems	11
2.1 Deterministic Inversion	11
2.2 Bayesian Inversion	13
Chapter 3. Software Framework	15
3.1 Firedrake	15
3.2 hIPPYfire	17
Chapter 4. Sample Problem	20
Chapter 5. Conclusion	28
Bibliography	29

List of Figures

3.1	Firedrake Abstractions	15
3.2	Firedrake v/s FEniCS flow	16
3.3	Flow and functionality of <code>hIPPYfire</code>	19
4.1	Variation of the true parameter and mean.	23
4.2	Variation of the true state and the observations.	24
4.3	Gradient and Hessian Checks obtained from <code>modelVerify</code>	25
4.4	Contours of the state and parameter	27

Chapter 1

Introduction

The advances made in the fields of high-performance computing have facilitated the development of large-scale solvers for *forward problems*. In a forward problem, inputs (such as the initial and boundary conditions, geometry, sources, material properties etc.) are supplied to the model of a physical system. The forward problem is then solved to determine the output quantities that are of interest. The ready availability of computational resources and the development of powerful discretization techniques that cater to a large variety of models have made the solutions of these forward problems extremely scalable. The vast amounts of data that are available, along with the subsequent improvements in data analysis techniques, have generated commensurate interest in extracting information about the physical model from the observed data. While this progress has been driven by machine learning algorithms, the estimation of physics-based models from data has been a crucial aspect of applied mathematics. Although significant research has already been conducted in this field [4, 23], the solution of inverse problems presents a different set of challenges.

Most PDE models of physical systems are characterized by a low-dimensional solution manifold. The inputs to these physical systems include initial and boundary conditions, geometry, sources, material and system properties, etc. The outputs, which are functions of some state variables, are obtained by solving the governing PDEs of the system. The low dimensionality of the map is attributed to the fact that the inputs and outputs are infinite-dimensional (or high-dimensional after discretization). The map from the inputs to the outputs must thus be smoothing, otherwise it may result in information loss. Inverse problems, on the other hand, involve estimation of inputs (commonly known as parameters) from noisy observations of a particular physical model (data). This low-dimensionality of the map causes the estimation of the parameters to be unstable in the presence of noise, which may be present in the system due to the observation process and/or model uncertainty. As a result, the problem is termed to be *ill-posed* due to the violation of one of the three conditions for well-posedness postulated by Hadamard [14]. This implies that despite the

high-dimensional nature of the data, it does not provide sufficient information to compute a unique solution for the parameter. A computational framework capable of solving such inverse problems, thus, became an important requirement.

This requirement was addressed by **hIPPYlib**, an inverse problem library that was capable of solving ill-posed large-scale deterministic and Bayesian inverse problems [27]. **hIPPYlib** poses the learning-from-data problem as an inverse problem governed by the forward PDE problem. It exploited the low-dimensionality of the parameter-to-observable map to estimate the components of the model at a cost that is measured in terms of the number of forward model solutions. **hIPPYlib** discretizes the forward map (which is one of the three components of an inverse problem, as discussed in Section 2) by using the functionality of the **FEniCS** library—a finite element library for partial differential equations (PDEs) [1]. Since the data structures used in the **FEniCS** library are wrappers on the PETSc library’s data structures, **hIPPYlib** uses PETSC [3] for its linear solvers and algebra operations. The **FEniCS** library contains different components with multiple layers of abstraction for ease of development and usage. However, given the mathematically complex nature of finite element problems, the consistent development of the **FEniCS** library and subsequently, that of **hIPPYlib**, requires individuals highly skilled in programming and FEM concepts.

Similar to the **FEniCS** project’s **DOLFIN** library [1, 16], **Firedrake** [18] is a library that provides automated solutions of partial differential equations (PDEs) using the finite element method (FEM). **Firedrake** managed to address the abovementioned problems by introducing a new layer of abstraction named **PyOP2** [19], while maintaining compatibility with **FEniCS** through the Unified Form Language (UFL) [2]. The **FEniCS** Project developed UFL [2] to express PDEs in a high-level mathematical language and seamlessly integrated it with their **DOLFIN** library. Well-defined domain-specific languages (DSLs) allow mathematicians to express their problem unambiguously, and **Firedrake** allows users to express finite element equations in the same syntax as that used in **DOLFIN** [16]. UFL is a symbolic language with well-defined finite element semantics, and its integration with **Firedrake** ensures that users of **FEniCS** can also use **Firedrake** with ease.

Thus, the need for an inverse problem library built on a modularized FEM solver like **Firedrake** led to the conception of **hIPPYfire**. This study presents the structure and implementation of **hIPPYfire**—an inverse problem library modelled similar to **hIPPYlib**, but built on **Firedrake**. The algorithms implemented in **hIPPYfire** are identical to those imple-

mented in `hIPPYlib`. However, there are minor differences in the utilities and linear algebra functions due to the currently limited functionality of the Firedrake library—all of which have been elaborated and expanded upon in the following sections of this report.

Section 2 provides a brief introduction to the theory behind inverse problems. The advantages of the Firedrake library over FEniCS are explained with some technical context in Section 3. Emphasis is also placed on the shortcomings of the Firedrake library and the development of custom linear algebra methods to address them. The structure of `hIPPYfire` is also discussed in Section 3. The `hIPPYfire` library is validated in Section 4 by creating a test case involving a Bayesian inverse problem. This section is followed by the conclusion (Section 5), which discusses the development roadmap for `hIPPYfire`.

Chapter 2

Theory Behind Inverse Problems

This section discusses two general approaches adopted to solve inverse problems—namely the deterministic framework [15, 30] and Bayesian framework [9, 22] for inverse problems. `hIPPYfire` currently implements an infinite-dimensional Bayesian framework [7], and consequently, additional emphasis has been placed on the theory of the latter framework.

Lower-case italic font is used to represent scalar valued functions like the parameter m , state u , observed data d etc. The discretized equivalents of these functions \mathbb{R}^n (where n is the discretization dimension) are denoted with a bold, lowercase font, such as \mathbf{m} , \mathbf{u} , \mathbf{d} etc. Vector functions in \mathbb{R}^2 or \mathbb{R}^3 are denoted in italic, bold, lowercase font, such as \mathbf{v} (velocity field). Infinite-dimensional spaces are represented through calligraphic font, such as \mathcal{A} . Scalars are denoted by using Greek font.

2.1 Deterministic Inversion

The solution of an inverse problem involves the inference of the parameter $m \in \mathcal{X}$, given data $\mathbf{d} \in \mathbb{R}^q$ by using the following *parameter-to-observable* map:

$$\mathcal{F}(m) = \mathbf{d} + \eta \tag{2.1.1}$$

Similar to *hIPPYlib*, the linear or non-linear *parameter-to-observable* (*p2o*) map is depicted as $\mathcal{F} : \mathcal{M} \rightarrow \mathbb{R}^q$. Note that $\mathcal{M} \subseteq L^2(\mathcal{D})$, where $\mathcal{D} \subset \mathbb{R}^d$ is a bounded domain. It is important to distinguish between the *p2o* map the forward problem. The forward problem is a map from the parameter m to the PDEs that govern the physical system, i.e., $m \rightarrow r(u, m)$, where $r(u, m)$ is the residual and $u \in \mathcal{V}$ is a state variable and \mathcal{V} is a Hilbert Space of functions defined on \mathcal{V} .

The forward problem is one of the three components of the forward map \mathcal{F} . The second component is the map $r : \mathcal{V} \times \mathcal{M} \rightarrow \mathcal{V}^*$. This map involves the solution of the governing PDEs. The third component is the observation operator (B), which maps u to the observable

$y \in \mathbb{R}^q$, which can be high or infinite dimensional. Thus, the map \mathcal{F} is now defined as:

$$\mathcal{F}(m) = \mathcal{B}(u), \text{ given, } r(u, m) = 0 \quad (2.1.2)$$

The noise, η , accounts for the difference between the observable y and data \mathbf{d} and is modelled as a Gaussian centred at 0 with a covariance Γ_{noise} . The source of the noise can be traced to imprecise measurements, model, and/or numerical errors. Although their exact values are not known, their statistical information (mean, variance, etc.), are known.

The ill-posedness of inverse problems [25], as discussed earlier, can primarily be attributed to their instability. The collapse of the spectrum of \mathcal{F} is the main culprit behind this phenomenon. The small eigenvalues of \mathcal{F} , which correspond to eigen function modes below the noise threshold, cause the noise in the data to blow up exponentially—rendering the inversion useless. The most obvious solution to this problem would be to discard the irrelevant eigenvalues. Although a truncated SVD presents a viable solution to the problem, the *Tikhonov Regularization* technique [13] allows us to formulate the inverse problem as an optimization problem instead of applying a filter. This can be mathematically represented as a non-linear least squares optimization problem:

$$\min_{m \in \mathcal{M}} \mathcal{J}(m) := \frac{1}{2} \|\mathcal{F}(m) - \mathbf{d}\|_{\Gamma_{noise}^{-1}}^2 + \mathcal{R}(m) \quad (2.1.3)$$

$\mathcal{J}(m)$ represents the cost functional. The first term on the RHS represents the misfit between $\mathcal{F}(m)$, weighted by the inverse noise covariance η_{noise}^{-1} , and data \mathbf{d} . The regularization parameter, $\mathcal{R}(m)$, ensures smoothness on the inversion parameter m .

Information on the first and second derivatives (gradient and Hessian, respectively) is required to solve the nonlinear optimization problem. Lagrangian techniques [26] were adopted in the development of the `hIPPYlib` library [27] to compute the actions of the Hessian and gradient. `hIPPYfire` makes use of the *Inexact Newton Conjugate Gradient Method* to solve the optimization problem, similar to the flow followed by `hIPPYlib` [27].

The regularization technique of solving inverse problems (2.1.1), although scalable with efficient inverse solvers, fails to account for uncertainties in the inferred parameters. This can be attributed to the fact that this approach only provides a point estimate of the inverse problem. Thus, it is not very useful for ill-posed problems with non-negligible noise—thereby prompting the conception of the Bayesian framework.

2.2 Bayesian Inversion

The Bayesian framework treats the inverse problem as statistical inference over a space of uncertain parameters. It computes a *posterior probability distribution* which represents the probability of the parameter being conditioned on the given data. One of the components of this framework is a *prior distribution*, which accounts for any constraints or assumptions on the data before data collection. This is combined with the *likelihood*, which expresses the probability of obtaining the observed data from a given set of parameters. This is achieved by analyzing and characterizing the posterior through mean estimation, sample drawing, analysis of the covariance etc.

However, complete characterization of the posterior is impractical for expensive PDE forward models, especially for ones in high dimensions that have been obtained after the discretization of infinite-dimensional parameter fields. Recently developed techniques, however, take advantage of the low dimensionality to address these problems—similar to the properties of the deterministic framework. Some of these include forward model reduction [12], Markov chain Monte-Carlo techniques that utilize the log-likelihood Hessian approximations [17], randomize-then-optimize techniques [31], etc. **hIPPYfire**, much like its predecessor [27], uses *Lagrangian approximation* of the posterior. This is extremely scalable and efficient—especially if the same properties that made the regularized Newton-CG method so powerful are used.

The infinite-dimensional Baye’s formula is given as:

$$\frac{d\mu_{post}}{d\mu_{prior}} \propto \pi_{like}(\mathbf{d}|m) \quad (2.2.1)$$

The LHS and RHS represent the Radon-Nikodym derivative [32] and likelihood, respectively. To study the impact of the noise on the likelihood, the noise is modelled as a centered Gaussian on \mathbb{R}^q with a covariance of Γ_{noise} . Thus, the likelihood can be expressed as

$$\pi_{like}(\mathbf{d}|m) \propto \exp(-\Phi(m)) \quad (2.2.2)$$

where $\Gamma(m) = \frac{1}{2} \|\mathcal{F}(m) - \mathbf{d}\|_{\Gamma_{noise}^{-1}}^2$ denotes the negative log-likelihood. The prior is chosen to be Gaussian, giving the following relation:

$$d\mu_{prior}(m) \propto \exp\left\{-\frac{1}{2} \|m - m_{pr}\|_{C_{prior}^{-1}}^2\right\} \quad (2.2.3)$$

where $m \sim \mathcal{N}(m_{pr}, C_{prior})$. For a parameter that represents a spatial field on $\mathcal{D} \in \mathbb{R}^d$, the prior covariance operator C_{prior} ensures that the computed parameter is smooth and

depends continuously on the data. The motivation behind this is to avoid the presence of discontinuous or rough components in the parameter, which makes it difficult to infer the latter from the data. The construction of the prior is similar to the methodology followed by `hIPPYlib` [27]. The equations of the likelihood (Eqn. (2.2.2)) and prior (Eqn.(2.2.3)) are substituted in the posterior distribution equation (Eqn. (2.2.1)) to yield:

$$d\mu_{post} \propto \exp\left\{-\frac{1}{2}\|\mathcal{F}(m) - \mathbf{d}\|_{\Gamma_{noise}^{-1}}^2 - \frac{1}{2}\|m - m_{pr}\|_{C_{prior}^{-1}}^2\right\} \quad (2.2.4)$$

The next step in the flow involves the computation of the MAP point m_{MAP} —which is the parameter field that maximizes the posterior distribution. It is obtained by solving the following optimization problem:

$$m_{MAP} := \underset{m \in \mathcal{M}}{\operatorname{argmin}} \frac{1}{2}\|\mathcal{F}(m) - \mathbf{d}\|_{\Gamma_{noise}^{-1}}^2 + \frac{1}{2}\|m - m_{pr}\|_{C_{prior}^{-1}}^2 \quad (2.2.5)$$

It is noteworthy that the prior performs the role of *Tikhonov Regularization* [13]. Similarities can be drawn between Eqn. (2.2.5) and Eqn. (2.1.3). In case the *p2o* map is non-linear, the posterior does not follow a Gaussian distribution. However, certain assumptions can be made on the noise covariance Γ_{noise} , number q of observations, and regularity of the *p2o* map \mathcal{F} , the Laplace approximation can be utilized to estimate the expected value of the prior [24, 10]. This is followed by the discretization of the Bayesian inverse problem, which has been explained in [7, 27]. The complete implementation of the Laplace approximation, along with the Bayesian discretization, in `hIPPYfire` is similar to that of `hIPPYlib` [27].

Chapter 3

Software Framework

3.1 Firedrake

As mentioned previously, one of the improvements Firedrake [18] implemented over FEniCS [1] is the creation of a new abstraction layer, namely PyOP2 [19], to distinguish between the local discretization or mathematical operators and their parallel execution over the mesh in the implementation layer. Firedrake models a finite problem as a combination of several abstractions—thereby allowing the user to follow a modularized approach while defining their problem. As mentioned earlier, Firedrake improved upon FEniCS by adopting a phi-

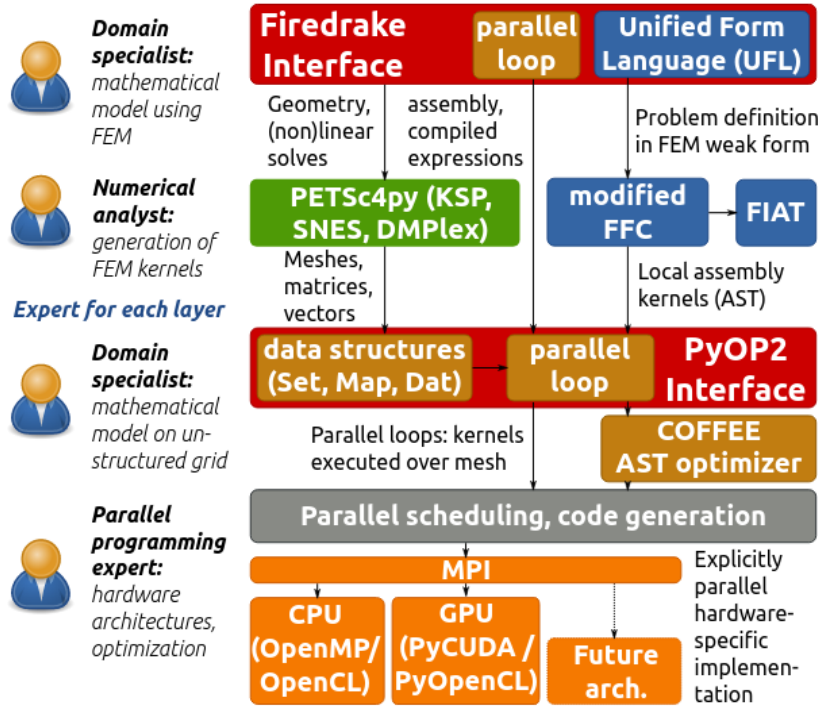


Figure 3.1: Firedrake Abstractions

Depiction of the separation of user concerns in Firedrake. Tools using FEniCS and PETSc are highlighted in blue and green respectively. The PyOP2 layers are shown in brown, while the backend engine is shown in orange [18].

losophy that emphasized on the separation of concerns—thereby providing a clear distinction between the mathematical and programming aspects of the library. Since a multidisciplinary

skillset, which ranges from mathematical expertise in numerical analyses to a deep understanding of parallel computation, is required for the development of these tools, it became more practical to develop abstract layers in the library that catered to a particular skillset. Firedrake introduced a new layer of abstraction named PyOP2 [16] that clearly formed a distinction between the finite element interface and the parallel execution of its algorithm over the mesh. PyOP2 thus creates a separation between the discretization of the mathematical operators and their parallel execution over the mesh. This has made the Firedrake codebase significantly more compact.

The cost of typical finite element problems can be attributed to data movement and floating point operations—both of which are proportional to the mesh size. These operations can be divided into two categories—custom mesh-defined data structure iterations and sparse linear algebra. While Firedrake utilizes PETSc [3] for the latter, PyOP2 was designed to address the former [19]. Additional information regarding the integration of PyOP2 with the Firedrake layer is shown in the figure below [18].

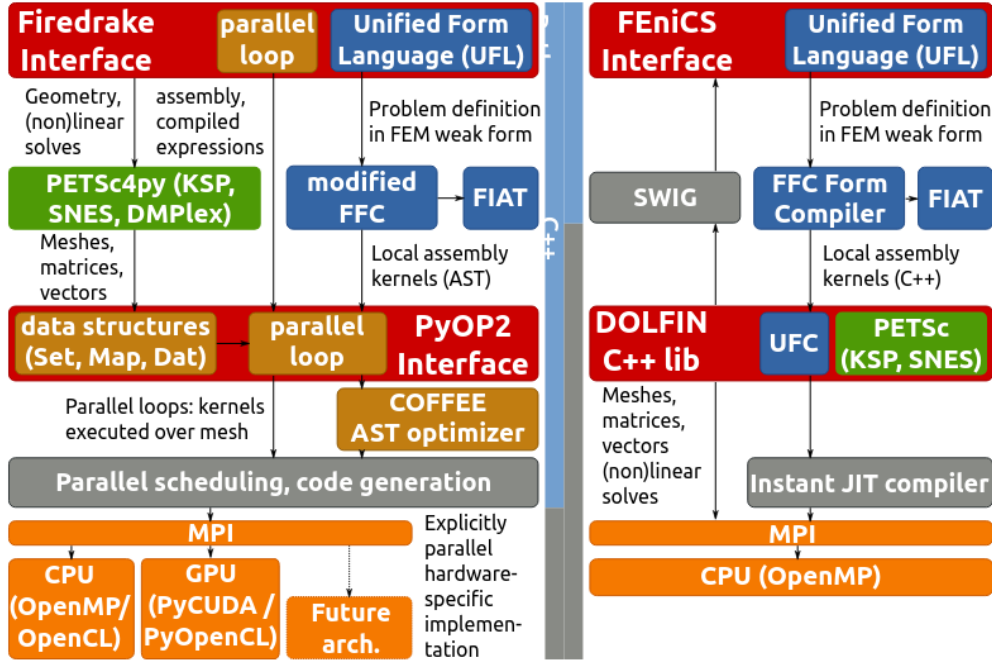


Figure 3.2: Firedrake v/s FEniCS flow

The clear distinction introduced in Firedrake is evident from the PyOP2 interface in its flow. [20].

The primary motivation behind using Firedrake can be attributed to its improved abstraction layer. However, a comparative analysis of the performances of Firedrake and FEniCS

was conducted [18]. Firedrake reported a better performance than FEniCS—however, no concrete reason has been provided for its superior performance.

3.2 hIPPYfire

hIPPYfire attempts to accomplish the same objective as that of **hIPPYlib**, i.e., implementation of scalable algorithms for PDE-based deterministic and Bayesian inverse problems. However, unlike its predecessor, it is built on Firedrake instead of FEniCS. The user is required to provide the PDE problem and likelihood in UFL [2], and **hIPPYfire** computes the gradient and Hessian. **hIPPYfire** is currently in development and does not support all the functionality of **hIPPYlib** at the timing of writing this report. Its different components have been summarized below:

- **Models:** The **modeling** module allows the user to specify information on the forward problem, misfit functional, and the prior.
 1. Forward Problem: This module computes the solutions of the forward, adjoint, and incremental problems. **hIPPYlib** and **hIPPYfire** both accept user input for the forward problem as a UFL form or user-defined object. The latter is required for transient inverse problems. However, if the forward problem is input as a UFL form, **hIPPYfire** computes the gradient and Hessian as well.
 2. Misfit: The misfit module evaluates the negative log-likelihood and its derivatives. Currently, the only misfit functional **hIPPYfire** provides support for is that of continuous observations—however, other functionals are currently in development.
 3. Prior: The prior computes the negative log-density and its derivatives, in addition to drawing samples and estimating the marginal variance. The user can select a Bilaplacian prior in **hIPPYfire**’s current implementation. There is a provision to accept user-defined priors as well.
 4. Model: The model is used to set up the *p2o* map. Its three components are computed from the abovementioned modules.
 5. Hessian: If the forward problem is input in a standard UFL form, **hIPPYfire** internally computes the Hessian of the forward map. The collapse of the spectrum

of the Hessian significantly influences the ill-posedness of the problem. However, the Hessian assumes a dense structure after discretization, thereby requiring forward and adjoint solves. Since the dimension of the Hessian is equal to that of the parameter, computing the Hessian for large-scale problems is not feasible. The rapidly decaying spectrum of the Hessian is exploited because the eigenvalues that tend to zero contain minimal information about the infinite-dimensional parameter field [11, 6].

In case of transient problems, the user will have to provide their own derivatives.

- **Algorithms:** The `algorithms` module contains an implementation of the inexact Newton-CG algorithm [5]. This is used to solve the deterministic inversion problem and compute the maximum *a posteriori* distribution (also called the MAP point) for Bayesian inversion problems. Following the computation of the gradient and Hessian forms, the Newton system is solved according to the Steihaug criterion [21]. Additional information regarding the implementation of this algorithm can be found in Villa et al. [29]. Custom linear algebra algorithms have also been implemented to perform basic matrix-vector operations that have been described later in this section.
- **Utilities:** The `utils` module contains certain helper functions to extract relevant data. The `vector2function` module wraps a discrete vector onto a continuous function, while the `rand` module generates random functions that are used to model the noise function and vector. The `parameterlist` module creates a custom list of parameters for the Newton solver.

One of the major challenges faced in the development of `hIPPYfire` was the limited functionality of Firedrake’s `firedrake.matrix.Matrix` data structure. Certain modules required matrix transpose and matrix-vector multiplication operations. Since these operations are not defined for the firedrake Matrix, Firedrake’s interface with PETSc was used to create linear algebra operations between Firedrake objects and their PETSc wrappers. These operations include matrix transpose (`Transpose()`), matrix-vector multiplication (`matVecMult()`), and matrix-transpose-vector multiplication (`matVecMultTranspose()`)—all of which have been defined in the `linalg` module. The `hIPPYlib` library provides APIs to create different kinds of solvers (the

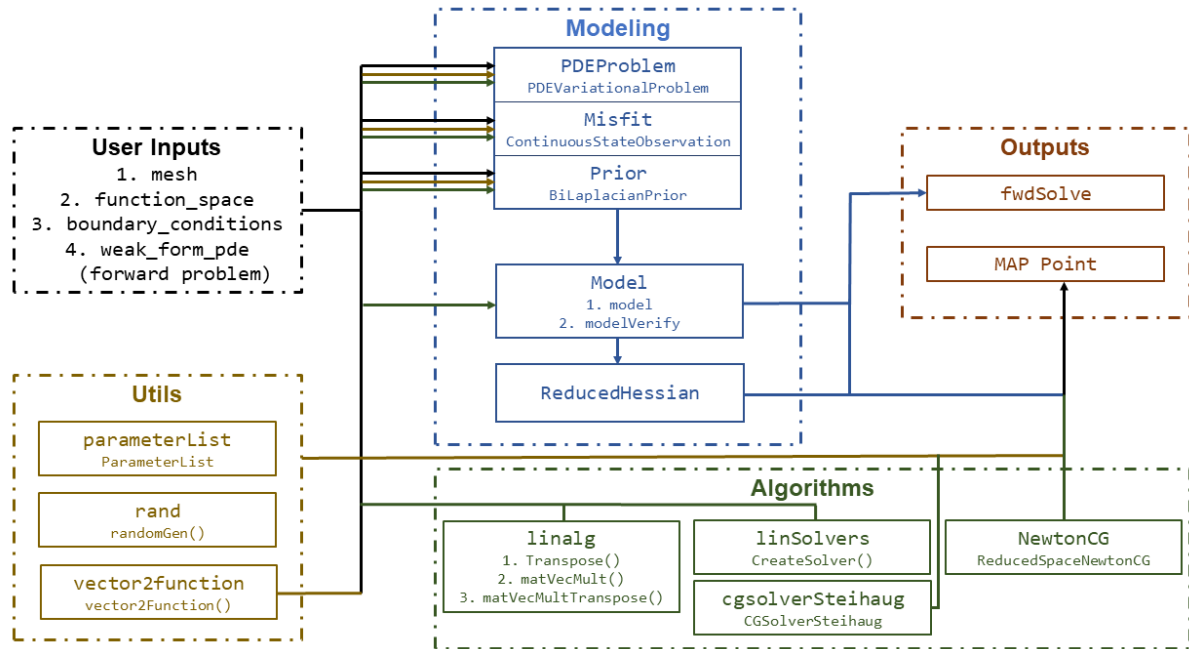


Figure 3.3: Flow and functionality of hIPPYfire

PETScKrylovSolver() and PETScLUSolver(), hIPPYfire provides a single API (CreateSolver()), which acts as a wrapper for Firedrake's linear solvers API and allows the user to create their custom solver.

Chapter 4

Sample Problem

In order to validate the different modules of `hIPPYfire`, one of `hIPPYlib`'s Bayesian inversion test cases, namely `3-SubsurfaceBayesian` [28], was recreated. This test case utilized all of `hIPPYfire`'s models that have been currently developed. Additional test cases shall be added for future functionality. The test case is briefly discussed below; a detailed description of the test case can be found in `hIPPYlib`'s repository [28].

The objective of this test case is to compute the uncertainty in the solution of an inverse problem governed by an elliptic PDE via the Bayesian inference framework. However, this test case only solves the forward problem and computes the MAP point, as described in Section 2. Please note that the following test case admits discretized expressions of the parameter space, i.e., they are finite-dimensional.

The objective of this sample problem is to compute the log coefficient field m for an elliptic PDE with continuous state observations. In addition to the mesh set-up parameters (such as the dimension, mesh divisions, function space, etc), the user is required to provide the weak form of the PDE and the corresponding boundary conditions for forward and adjoint problems. The weak form of the PDE is given as follows:

$$\begin{aligned} e^m \langle \nabla u, \nabla p \rangle &= \langle f, p \rangle \text{ in } \Omega, \\ u(\mathbf{x}, \mathbf{y}) &= \mathbf{y} \text{ on } \Gamma_D = \Gamma_{bottom} \bigcup \Gamma_{top} \end{aligned} \tag{4.0.1}$$

where $\Omega = (0, 1) \times (0, 1)$ and Γ_D is the union of the top and bottom boundaries of Ω . u is the state variable, $f \in \mathcal{L}^2(\Omega)$ is a source term (that is equated to zero in the current case), and $\langle \cdot, \cdot \rangle$ represents the standard inner product on $\mathcal{L}^2(\Omega)$. The spaces V_y and V_0 are defined as follows:

$$\begin{aligned} V_y &= \left\{ v \in H^1(\Omega) \text{ s.t. } v = \mathbf{y} \text{ on } \Gamma_D \right\} \\ V_0 &= \left\{ v \in H^1(\Omega) \text{ s.t. } v = 0 \text{ on } \Gamma_D \right\} \end{aligned}$$

where $H^1(\Omega)$ is the Sobolev space whose derivatives are in $\mathcal{L}^2(\Omega)$. In the weak form of the PDE (Eqn. 4.0.1), we attempt to find $u \in V_y$ using the test function $p \in V_0$.

The prior is a Gaussian distribution $\mathcal{N}(m_{pr}, C_{prior})$ s.t. $C_{prior} = \mathcal{A}^{-2}$, as shown in [22]. \mathcal{A} is a differential operator with a domain $H^1(\Omega)$ and action as given below:

$$\left\{ \begin{array}{ll} -\gamma \nabla \cdot (\Theta \nabla m) + \delta m & \text{in } \Omega \\ \Theta \nabla m \cdot \mathbf{n} + \beta m & \text{in } \partial\Omega \end{array} \right\} \quad (4.0.2)$$

$\beta \propto \sqrt{\gamma\delta}$ is the optimal robin coefficient [8] to minimize boundary artifacts. Θ is an s.p.d anisotropic tensor of the form:

$$\left\{ \begin{array}{ll} \theta_1 \sin(\alpha)^2 & (\theta_1 - \theta_2) \sin(\alpha) \cos(\alpha) \\ (\theta_1 - \theta_2) \sin(\alpha) \cos(\alpha) & \theta_2 \cos(\alpha)^2 \end{array} \right\} \quad (4.0.3)$$

The prior is chosen such that the well-posedness of this problem is ensured.

The log-likelihood (data misfit) functional is specified next. The noisy continuous data observations, \mathbf{d} , are approximated s.t. $\mathbf{d} \in \mathbb{R}^q$. They represent the observations of the state u at $q = 16641$ random locations uniformly distributed in Ω

@Umberto, kindly request you to verify the above statement. This is the reason why we have points in the RHS of the observations graph, even though it is a continuous state observation. `misfit.d` is initialized as `pde.generate_state()` vector. If my explanation above is incorrect, could you please assist me in rectifying it? Thank you.

$$d = \mathcal{B}u + \eta \quad (4.0.4)$$

where $\mathcal{B} : V_y \rightarrow \mathbb{R}^q$ is a linear observation operator. η represents the noise vector and is a multivariate Gaussian variable with a distribution $\mathcal{N}(0, \Gamma_{noise})$ s.t. $\Gamma_{noise} = \sigma^2 \mathbf{I}$, where $\sigma = 0.01$ and $\mathbf{I} = \mathbb{R}^{q \times q}$

The likelihood is then computed as follows:

$$\pi_{like}(\mathbf{d}|\mathbf{m}) = \exp\left(-\frac{1}{2}(\mathcal{B}u(m) - \mathbf{d})^T \Gamma_{noise}^{-1} (\mathcal{B}u(m) - \mathbf{d})\right) \quad (4.0.5)$$

where m_{pr} is the prior mean of the log coefficient field m . The above equation presents a discretized version of the likelihood function. The mean of the posterior distribution, also referred to as m_{MAP} , is the parameter vector that maximizes the posterior. It is computed by solving the following variational non-linear least squares optimization problem:

$$\mathbf{m}_{MAP} := \underset{m \in H^1(\Omega)}{\operatorname{argmin}} \mathcal{J}(\mathbf{m}) := \left(\frac{1}{2} \|\mathcal{B}u(m) - \mathbf{d}\|_{\Gamma_{noise}^{-1}}^2 + \frac{1}{2} \|\mathbf{m} - \mathbf{m}_{prior}\|_{C_{prior}^{-1}}^2\right) \quad (4.0.6)$$

This optimization problem has been solved using the inexact Newton-CG algorithm, which has been described in Appendix A of Villa et al. [28].

@Umberto, Since we do not go beyond the calculation of the MAP point, I have not included any more equations. Would you like me to include the Gradient and Hessian actions of the cost functional as well? Please let me know, and I shall include them appropriately. Although I might need your help in writing those since we only have the weak form of the PDE.

The problem flow demonstrated above is represented through a few code snippets below.

- **Mesh and FEM setup:** A two-dimensional unit square mesh is created with a P2 finite element space for `state` and `adjoint` variables and P1 for `parameter`.

```

1     ndim = 2
2     nx = 64
3     ny = 64
4     mesh = fd.UnitSquareMesh(nx, ny)
5     Vh2 = fd.FunctionSpace(mesh, 'Lagrange', 2)
6     Vh1 = fd.FunctionSpace(mesh, 'Lagrange', 1)
7     Vh = [Vh2, Vh1, Vh2]
8

```

- **Forward Problem:** As mentioned in Section 3, the `PDEVariationalProblem` class sets up the forward problem component of the *p2o* map. In addition to the finite element components defined above, it requires an expression of the weak form of the PDE (given by `pde_varf`) and boundary conditions for the forward (`bc`) and incremental and adjoint problems (`bc0`). The `PDEVariationalProblem` class solves the forward/adjoint and incremental problems and computes the relevant partial derivatives with respect to the state, parameter, and adjoint variables.

```

1     u_bdr = fd.SpatialCoordinate(mesh)[1]
2     u_bdr0 = fd.Constant(0.0)
3
4     bc = fd.DirichletBC(Vh[STATE], u_bdr, [3, 4]) # [3, 4]
5     indicates that bc is applied to y == 0 amd y ==1
6     bc0 = fd.DirichletBC(Vh[STATE], u_bdr0, [3, 4])
7
8     f = fd.Constant(1.0)
9
10    def pde_varf(u, m, p):
11        return ufl.exp(m) * ufl.inner(ufl.grad(u), ufl.grad(p)) *
12        ufl.dx - f * p * ufl.dx
13
14    pde = PDEVariationalProblem(Vh, pde_varf, bc, bc0,
15                                is_fwd_linear=True)

```

The `is_fwd_linear=True` flag allows the user to set a non-linear forward map as well.

- **Prior setup:** The class `BiLaplacianPrior` creates a Gaussian prior with zero average, Additional information regarding the covariance can be found in Villa et al. [28].

```

1      pr = BiLaplacianPrior(Vh[PARAMETER], gamma, delta, robin_bc=
      True)
2      x = fd.SpatialCoordinate(mesh)
3      mtrue = fd.interpolate(fd.sin(x[0])*fd.cos(x[1]), Vh[
      PARAMETER]).vector()
4      m0 = fd.interpolate(fd.sin(x[0]), Vh[PARAMETER]).vector()
5      objs = [fd.Function(Vh[PARAMETER], mtrue), fd.Function(Vh[
      PARAMETER], pr.mean)]
6

```

The true parameter, `mtrue` is initialized to be a known analytic function to validate the accuracy of the parameters computed by `hIPPYfire`. Plots of the vectors `mtrue` and `pr.mean` are generated and shown below. For the purpose of validation through this test case, `mtrue` is a known function and not randomly generated, as in `hIPPYlib` [28]




Figure 4.1: Variation of the true parameter and mean.

- **Misfit:** hIPPYfire currently provides support for `ContinuousStateObservation`, which sets up the observation parameter \mathcal{B} . The observables which shall provide our input data are first generated by solving the forward problem by using the true parameter \mathbf{m}_{true}

```

1      misfit = ContinuousStateObservation(Vh[STATE], ufl.dx, bcs=
      bc0)
2      misfit.noise_variance = 1e-4
3      utrue = pde.generate_state()
4      x = [utrue, mtrue.vector(), None]
5      pde.solveFwd(x[STATE], x)
6      misfit.d.axy(1., utrue)
7      misfit.d.axy(float(np.sqrt(misfit.noise_variance)),
      randomGen(Vh[STATE]).vector())
8

```

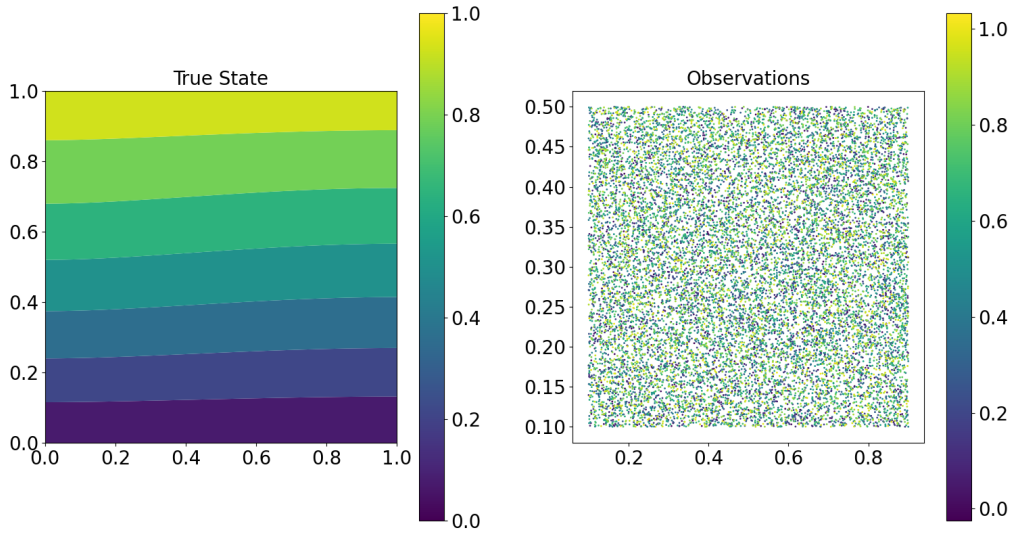


Figure 4.2: Variation of the true state and the observations.

- **Model:** The model, which is created by the `model` class, depends on three components—namely `PDEVariationalProblem`, `misfit`, and `prior`. The `PDEVariationalProblem` provides solutions for the forward and adjoint problems and incremental forward and adjoint problems. The `prior` applies the regularization operator to a vector, while the `misfit` computes the cost functional and partial derivatives with respect to the state

and parameter variables. Forward finite differences are used to test the model through the `modelVerify` module.

```
1 model = Model(pde, pr, misfit)
2 eps, err_grad, err_H = modelVerify(model, m0, misfit_only=
3 False)
```

```
1 (yy, H xx) - (xx, H yy) = 0.0
2
```




Figure 4.3: Gradient and Hessian Checks obtained from `modelVerify`

@Umberto, I am currently working to fix the aspect ratio of the gradient check plots.
Should be fixed by EOD

- **MAP Point:** The Newton-CG method is used to compute the MAP point.

```
1 m = pr.mean.copy()
2 solver = ReducedSpaceNewtonCG(model)
```

```

3 solver.parameters["rel_tolerance"] = 1e-6
4 solver.parameters["abs_tolerance"] = 1e-12
5 solver.parameters["max_iter"] = 25
6 solver.parameters["GN_iter"] = 5
7 solver.parameters["globalization"] = "LS"
8 solver.parameters["LS"]["c_armijo"] = 1e-4
9 x = solver.solve([None, m, None])
10

```

The following output was obtained:

```

1 Relative/Absolute residual less than tol
2 Converged in 19 iterations with final norm 5.65259722013104e-08
3
4 It cg_it cost misfit reg (g,dm)
5 1 1 9.452277e-01 8.990612e-01 4.616652e-02 -2.456844e+01
6 2 2 4.228824e-01 3.488744e-01 7.400803e-02 -1.043988e+00
7 3 5 4.027649e-01 3.196818e-01 8.308311e-02 -4.042007e-02
8 4 7 4.026241e-01 3.196381e-01 8.298594e-02 -3.009693e-04
9 5 7 4.026228e-01 3.196101e-01 8.301271e-02 -2.859017e-06
10 6 8 4.026228e-01 3.196134e-01 8.300935e-02 -4.118272e-08
11 8.007385e-04 1.000000e+00 2.179740e-03
12 5.941225051879883 Execution time
13 Converged in 6 iterations.
14 Termination reason: Norm of the gradient less than tolerance
15 Final gradient norm: 7.784347671669573e-06
16 Final cost: 0.40262278305284716
17

```

A mesh independence study was conducted to establish that the convergence of the problem is observed for different sizes. The current sample assumes a uniform square mesh with 64 divisions. Experiments were conducted for mesh sizes of 128 and 256 divisions. Convergence was reported for 6–7 iterations for all mesh sizes.




Figure 4.4: Contours of the state and parameter

Chapter 5

Conclusion

This report presents the development of a software library named **hIPPYfire** to solve large-scale deterministic and Bayesian inverse problems that are governed by partial differential equations. Its structure and algorithms are identical to those implemented in **hIPPYlib** [27] with a similar underlying principal of exploiting the effective low-dimensionality of the parameter-to-observable map to develop scalable algorithms. The equations describing the computation of the forward map and MAP point have been described in Sections 2.1 and 2.2. However, unlike **hIPPYlib**, **hIPPYfire** uses Firedrake [20] for the finite element solution of the PDE problem. Firedrake improves upon FEniCS by providing an additional layer of abstraction named PyOP2 [19], thereby instituting a clear distinction between the mathematical and parallel execution aspects of the algorithm (Section 3.1. This encourages steady and consistent development of Firedrake, and consequently, **hIPPYfire**. A detailed description of the three components of the *p2o* map and their implementation has been provided in Section 3.2. The functionality of the library is validated by solving the forward problem and computing the MAP point for a Bayesian inversion test case in Section 4.

The promising results provide sufficient incentives to further develop **hIPPYfire**. Functionality to compute the low-rank approximation of the Hessian misfit at the MAP point and the prior and Laplace approximations of the posterior pointwise variance fields remain to be added. However, the consistent addition of new functionality to the Firedrake library provides plenty of motivation to further develop the **hIPPYfire** library.

Bibliography

- [1] Martin Alnæs, Jan Blechta, Johan Hake, August Johansson, Benjamin Kehlet, Anders Logg, Chris Richardson, Johannes Ring, Marie E Rognes, and Garth N Wells. The fenics project version 1.5. *Archive of Numerical Software*, 3(100), 2015.
- [2] Martin S Alnæs, Anders Logg, Kristian B Ølgaard, Marie E Rognes, and Garth N Wells. Unified form language: A domain-specific language for weak formulations of partial differential equations. *ACM Transactions on Mathematical Software (TOMS)*, 40(2):1–37, 2014.
- [3] Satish Balay, Kris Buschelman, William D Gropp, Dinesh Kaushik, Matthew G Knepley, L Curfman McInnes, Barry F Smith, and Hong Zhang. Petsc. See <http://www.mcs.anl.gov/petsc>, 2001.
- [4] H Thomas Banks and Karl Kunisch. *Estimation techniques for distributed parameter systems*. Springer Science & Business Media, 2012.
- [5] Alfio Borzì and Volker Schulz. *Computational optimization of systems governed by partial differential equations*. SIAM, 2011.
- [6] Tan Bui-Thanh and Omar Ghattas. Analysis of the hessian for inverse scattering problems: I. inverse shape scattering of acoustic waves. *Inverse Problems*, 28(5):055001, 2012.
- [7] Tan Bui-Thanh, Omar Ghattas, James Martin, and Georg Stadler. A computational framework for infinite-dimensional bayesian inverse problems part i: The linearized case, with application to global seismic inversion. *SIAM Journal on Scientific Computing*, 35(6):A2494–A2523, 2013.
- [8] Yair Daon and Georg Stadler. Mitigating the influence of the boundary on pde-based covariance operators. *arXiv preprint arXiv:1610.05280*, 2016.
- [9] Masoumeh Dashti and Andrew M Stuart. The bayesian approach to inverse problems. In *Handbook of uncertainty quantification*, pages 311–428. Springer, 2017.

- [10] Michael Evans and Timothy Swartz. *Approximating integrals via Monte Carlo and deterministic methods*, volume 20. OUP Oxford, 2000.
- [11] H Pearl Flath, Lucas C Wilcox, Volkan Akçelik, Judith Hill, Bart van Bloemen Waanders, and Omar Ghattas. Fast algorithms for bayesian uncertainty quantification in large-scale linear inverse problems based on low-rank partial hessian approximations. *SIAM Journal on Scientific Computing*, 33(1):407–432, 2011.
- [12] David Galbally, Krzysztof Fidkowski, Karen Willcox, and Omar Ghattas. Non-linear model reduction for uncertainty quantification in large-scale inverse problems. *International journal for numerical methods in engineering*, 81(12):1581–1608, 2010.
- [13] Gene H Golub, Per Christian Hansen, and Dianne P O’Leary. Tikhonov regularization and total least squares. *SIAM journal on matrix analysis and applications*, 21(1):185–194, 1999.
- [14] Jacques Hadamard. *Lectures on Cauchy’s problem in linear partial differential equations*, volume 15. Yale university press, 1923.
- [15] Eric Laloy, Niklas Linde, Cyprien Ruffino, Romain Hérault, Gilles Gasso, and Diederik Jacques. Gradient-based deterministic inversion of geophysical data with generative adversarial networks: is it feasible? *Computers & Geosciences*, 133:104333, 2019.
- [16] Anders Logg and Garth N Wells. Dolfin: Automated finite element computing. *ACM Transactions on Mathematical Software (TOMS)*, 37(2):1–28, 2010.
- [17] Noemi Petra, Hongyu Zhu, Georg Stadler, Thomas JR Hughes, and Omar Ghattas. An inexact gauss-newton method for inversion of basal sliding and rheology parameters in a nonlinear stokes ice sheet model. *Journal of Glaciology*, 58(211):889–903, 2012.
- [18] Florian Rathgeber, David A Ham, Lawrence Mitchell, Michael Lange, Fabio Luporini, Andrew TT McRae, Gheorghe-Teodor Bercea, Graham R Markall, and Paul HJ Kelly. Firedrake: automating the finite element method by composing abstractions. *ACM Transactions on Mathematical Software (TOMS)*, 43(3):1–27, 2016.

- [19] Florian Rathgeber, Graham R Markall, Lawrence Mitchell, Nicolas Lorient, David A Ham, Carlo Bertolli, and Paul HJ Kelly. Pyop2: A high-level framework for performance-portable simulations on unstructured meshes. In *2012 SC Companion: High Performance Computing, Networking Storage and Analysis*, pages 1116–1123. IEEE, 2012.
- [20] Florian Rathgeber, Lawrence Mitchell, David Ham, Michael Lange, Andrew McRae, Fabio Luporini, Gheorghe-teodor Bercea, and Paul Kelly. Firedrake: Re-imagining fenics by composing domainspecific abstractions, 2014.
- [21] Trond Steihaug. Local and superlinear convergence for truncated iterated projections methods. *Mathematical Programming*, 27:176–190, 1983.
- [22] Andrew M Stuart. Inverse problems: a bayesian perspective. *Acta numerica*, 19:451–559, 2010.
- [23] Timothy John Sullivan. *Introduction to uncertainty quantification*, volume 63. Springer, 2015.
- [24] Luke Tierney and Joseph B Kadane. Accurate approximations for posterior moments and marginal densities. *Journal of the american statistical association*, 81(393):82–86, 1986.
- [25] Andrei Nikolaevich Tikhonov. On the solution of ill-posed problems and the method of regularization. In *Doklady akademii nauk*, volume 151, pages 501–504. Russian Academy of Sciences, 1963.
- [26] Fredi Tröltzsch. *Optimal control of partial differential equations: theory, methods, and applications*, volume 112. American Mathematical Soc., 2010.
- [27] Umberto Villa, Noemi Petra, and Omar Ghattas. hippylib: An extensible software framework for large-scale inverse problems. *Journal of Open Source Software*, 3(30), 2018.
- [28] Umberto Villa, Noemi Petra, and Omar Ghattas. hIPPYlib: An Extensible Software Framework for Large-Scale Inverse Problems Governed by PDEs, 2 2020.
- [29] Umberto Villa, Noemi Petra, and Omar Ghattas. hIPPYlib User Manual: Version 2. 6 2020.

- [30] Curtis R. Vogel. *Computational Methods for Inverse Problems*. Society for Industrial and Applied Mathematics, 2002.
- [31] Kainan Wang, Tan Bui-Thanh, and Omar Ghattas. A randomized maximum a posteriori method for posterior sampling of high dimensional nonlinear bayesian inverse problems. *SIAM Journal on Scientific Computing*, 40(1):A142–A171, 2018.
- [32] David Williams. *Probability with martingales*. Cambridge university press, 1991.